

```
import matplotlib.pyplot as plt
import os
import re
import shutil
import string
import tensorflow as tf

from tensorflow.keras import layers
from tensorflow.keras import losses

print(tf.__version__)

url = "https://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz"

dataset = tf.keras.utils.get_file("aclImdb_v1", url,
                                   untar=True, cache_dir='.',
                                   cache_subdir='')

dataset_dir = os.path.join(os.path.dirname(dataset), 'aclImdb')

os.listdir(dataset_dir)

train_dir = os.path.join(dataset_dir, 'train')
os.listdir(train_dir)

sample_file = os.path.join(train_dir, 'pos/1181_9.txt')
with open(sample_file) as f:
    print(f.read())

remove_dir = os.path.join(train_dir, 'unsup')
shutil.rmtree(remove_dir)

batch_size = 32
seed = 42

raw_train_ds = tf.keras.utils.text_dataset_from_directory(
    'aclImdb/train',
    batch_size=batch_size,
    validation_split=0.2,
    subset='training',
    seed=seed)

for text_batch, label_batch in raw_train_ds.take(1):
    for i in range(3):
        print("Review", text_batch.numpy()[i])
        print("Label", label_batch.numpy()[i])

print("Label 0 corresponds to", raw_train_ds.class_names[0])
print("Label 1 corresponds to", raw_train_ds.class_names[1])

raw_val_ds = tf.keras.utils.text_dataset_from_directory(
    'aclImdb/train',
    batch_size=batch_size,
    validation_split=0.2,
    subset='validation',
    seed=seed)

raw_test_ds = tf.keras.utils.text_dataset_from_directory(
    'aclImdb/test',
    batch_size=batch_size)

def custom_standardization(input_data):
    lowercase = tf.strings.lower(input_data)
    stripped_html = tf.strings.regex_replace(lowercase, '<br />', ' ')
    return tf.strings.regex_replace(stripped_html,
                                     '[%s]' % re.escape(string.punctuation),
                                     '')

max_features = 10000
sequence_length = 250

vectorize_layer = layers.TextVectorization(
    standardize=custom_standardization,
    max_tokens=max_features,
    output_mode='int',
    output_sequence_length=sequence_length)
```

```

train_text = raw_train_ds.map(lambda x, y: x)
vectorize_layer.adapt(train_text)

def vectorize_text(text, label):
    text = tf.expand_dims(text, -1)
    return vectorize_layer(text), label

text_batch, label_batch = next(iter(raw_train_ds))
first_review, first_label = text_batch[0], label_batch[0]
print("Review", first_review)
print("Label", raw_train_ds.class_names[first_label])
print("Vectorized review", vectorize_text(first_review, first_label))

print("1287 ---> ", vectorize_layer.get_vocabulary()[1287])
print(" 313 ---> ", vectorize_layer.get_vocabulary()[313])
print('Vocabulary size: {}'.format(len(vectorize_layer.get_vocabulary()))))

train_ds = raw_train_ds.map(vectorize_text)
val_ds = raw_val_ds.map(vectorize_text)
test_ds = raw_test_ds.map(vectorize_text)

AUTOTUNE = tf.data.AUTOTUNE

train_ds = train_ds.cache().prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
test_ds = test_ds.cache().prefetch(buffer_size=AUTOTUNE)

embedding_dim = 16
model = tf.keras.Sequential([
    layers.Embedding(max_features, embedding_dim),
    layers.Dropout(0.2),
    layers.GlobalAveragePooling1D(),
    layers.Dropout(0.2),
    layers.Dense(1, activation='sigmoid')])

model.summary()

model.compile(loss=losses.BinaryCrossentropy(),
              optimizer='adam',
              metrics=[tf.metrics.BinaryAccuracy(threshold=0.5)])

epochs = 10
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs)

loss, accuracy = model.evaluate(test_ds)

print("Loss: ", loss)
print("Accuracy: ", accuracy)

history_dict = history.history
history_dict.keys()

acc = history_dict['binary_accuracy']
val_acc = history_dict['val_binary_accuracy']
loss = history_dict['loss']
val_loss = history_dict['val_loss']

epochs = range(1, len(acc) + 1)

plt.plot(epochs, loss, 'bo', label='Training loss')

plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')

```

```
plt.xlabel('epochs')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')

plt.show()

export_model = tf.keras.Sequential([
    vectorize_layer,
    model,
    layers.Activation('sigmoid')
])

export_model.compile(
    loss=losses.BinaryCrossentropy(from_logits=False), optimizer="adam", metrics=['accuracy']
)

metrics = export_model.evaluate(raw_test_ds, return_dict=True)
print(metrics)

examples = tf.constant([
    "The movie was great!",
    "The movie was okay.",
    "The movie was terrible..."
])

export_model.predict(examples)
```

```

2.17.1
Downloading data from https://ai.stanford.edu/~amaas/data/sentiment/aclimdb\_v1.tar.gz
84125825/84125825 ————— 6s 0us/step
Rachel Griffiths writes and directs this award winning short film. A heartwarming story about coping with grief and cherishing the memor
Found 25000 files belonging to 2 classes.
Using 20000 files for training.
Review b'"Pandemonium" is a horror movie spoof that comes off more stupid than funny. Believe me when I tell you, I love comedies. Espec
Label 0
Review b'"David Mamet is a very interesting and a very un-equal director. His first movie 'House of Games' was the one I liked best, and
Label 0
Review b'"Great documentary about the lives of NY firefighters during the worst terrorist attack of all time.. That reason alone is why t
Label 1
Label 0 corresponds to neg
Label 1 corresponds to pos
Found 25000 files belonging to 2 classes.
Using 5000 files for validation.
Found 25000 files belonging to 2 classes.
Review tf.Tensor(b'"Silent Night, Deadly Night 5 is the very last of the series, and like part 4, it\'s unrelated to the first three exce
Label neg
Vectorized review (<tf.Tensor: shape=(1, 250), dtype=int64, numpy=
array([[1287, 313, 2380, 313, 661, 7, 2, 52, 229, 5, 2,
      200, 3, 38, 170, 669, 29, 5492, 6, 2, 83, 297,
      549, 32, 410, 3, 2, 186, 12, 29, 4, 1, 191,
      510, 549, 6, 2, 8229, 212, 46, 576, 175, 168, 20,
      1, 5361, 290, 4, 1, 761, 969, 1, 3, 24, 935,
      2271, 393, 7, 1, 1675, 4, 3747, 250, 148, 4, 112,
      436, 761, 3529, 548, 4, 3633, 31, 2, 1331, 28, 2096,
      3, 2912, 9, 6, 163, 4, 1006, 20, 2, 1, 15,
      85, 53, 147, 9, 292, 89, 959, 2314, 984, 27, 762,
      6, 959, 9, 564, 18, 7, 2140, 32, 24, 1254, 36,
      1, 85, 3, 3298, 85, 6, 1410, 3, 1936, 2, 3408,
      301, 965, 7, 4, 112, 740, 1977, 12, 1, 2014, 2772,
      3, 4, 428, 3, 5177, 6, 512, 1254, 1, 278, 27,
      139, 25, 308, 1, 579, 5, 259, 3529, 7, 92, 8981,
      32, 2, 3842, 230, 27, 289, 9, 35, 2, 5712, 18,
      27, 144, 2166, 56, 6, 26, 46, 466, 2014, 27, 40,
      2745, 657, 212, 4, 1376, 3002, 7080, 183, 36, 180, 52,
      920, 8, 2, 4028, 12, 969, 1, 158, 71, 53, 67,
      85, 2754, 4, 734, 51, 1, 1611, 294, 85, 6, 2,
      1164, 6, 163, 4, 3408, 15, 85, 6, 717, 85, 44,
      5, 24, 7158, 3, 48, 604, 7, 11, 225, 384, 73,
      65, 21, 242, 18, 27, 120, 295, 6, 26, 667, 129,
      4028, 948, 6, 67, 48, 158, 93, 1]])>, <tf.Tensor: shape=(), dtype=int32, numpy=0>)
1287 ---> silent
313 ---> night
Vocabulary size: 10000
Model: "sequential"

```

Layer (type)	Output Shape	Param #
embedding ( <a href="#">Embedding</a> )	?	0 (unbuilt)
dropout ( <a href="#">Dropout</a> )	?	0 (unbuilt)
global_average_pooling1d ( <a href="#">GlobalAveragePooling1D</a> )	?	0 (unbuilt)
dropout_1 ( <a href="#">Dropout</a> )	?	0 (unbuilt)
dense ( <a href="#">Dense</a> )	?	0 (unbuilt)

Total params: 0 (0.00 B)  
 Trainable params: 0 (0.00 B)  
 Non-trainable params: 0 (0.00 B)

```

Epoch 1/10
625/625 ————— 7s 10ms/step - binary_accuracy: 0.5813 - loss: 0.6805 - val_binary_accuracy: 0.7398 - val_loss: 0.6076
Epoch 2/10
625/625 ————— 8s 6ms/step - binary_accuracy: 0.7619 - loss: 0.5742 - val_binary_accuracy: 0.8152 - val_loss: 0.4931
Epoch 3/10
625/625 ————— 5s 6ms/step - binary_accuracy: 0.8287 - loss: 0.4617 - val_binary_accuracy: 0.8308 - val_loss: 0.4247
Epoch 4/10
625/625 ————— 6s 7ms/step - binary_accuracy: 0.8529 - loss: 0.3907 - val_binary_accuracy: 0.8432 - val_loss: 0.3841
Epoch 5/10
625/625 ————— 6s 9ms/step - binary_accuracy: 0.8691 - loss: 0.3457 - val_binary_accuracy: 0.8460 - val_loss: 0.3608
Epoch 6/10
625/625 ————— 5s 9ms/step - binary_accuracy: 0.8828 - loss: 0.3133 - val_binary_accuracy: 0.8500 - val_loss: 0.3449
Epoch 7/10
625/625 ————— 9s 6ms/step - binary_accuracy: 0.8911 - loss: 0.2885 - val_binary_accuracy: 0.8550 - val_loss: 0.3326
Epoch 8/10
625/625 ————— 6s 8ms/step - binary_accuracy: 0.9007 - loss: 0.2678 - val_binary_accuracy: 0.8590 - val_loss: 0.3217
Epoch 9/10
625/625 ————— 10s 7ms/step - binary_accuracy: 0.9074 - loss: 0.2522 - val_binary_accuracy: 0.8582 - val_loss: 0.3182
Epoch 10/10
625/625 ————— 5s 7ms/step - binary_accuracy: 0.9137 - loss: 0.2378 - val_binary_accuracy: 0.8606 - val_loss: 0.3140
787/787 ————— 4s 4ms/step - binary_accuracy: 0.8541 - loss: 0.3345

```

Loss: 0.33228880167007446  
Accuracy: 0.854200005531311

